

```
In [1]: import pandas as pd
        from sklearn.model_selection import train_test_split
        import random
        import numpy as np
        import seaborn as sns
        import matplotlib.pyplot as plt
```

```
In [2]: bank = pd.read_csv("C:/Users/sheri/Desktop/projectbank.csv")
```

```
In [3]: bank.shape
```

```
Out[3]: (45211, 17)
```

```
In [4]: bank.head()
```

```
Out[4]:
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may	

```
In [5]: bank.isna().sum()
```

```
Out[5]: age          0
        job          0
        marital      0
        education    0
        default      0
        balance      0
        housing      0
        loan         0
        contact      0
        day          0
        month        0
        duration     0
        campaign     0
        pdays        0
        previous     0
        poutcome     0
        y            0
        dtype: int64
```

```
In [6]: bank.dtypes
        #There are numerical and categorical data.
```

```
Out[6]: age          int64
```

```

job            object
marital        object
education      object
default        object
balance        int64
housing        object
loan           object
contact        object
day            int64
month          object
duration       int64
campaign       int64
pdays        int64
previous       int64
poutcome      object
y              object
dtype: object

```

```
In [7]: bank.columns
```

```
Out[7]: Index(['age', 'job', 'marital', 'education', 'default', 'balance', 'housing',
              'loan', 'contact', 'day', 'month', 'duration', 'campaign', 'pdays',
              'previous', 'poutcome', 'y'],
              dtype='object')
```

```
In [8]: bank.median()
```

```
Out[8]: age            39.0
balance        448.0
day            16.0
duration       180.0
campaign        2.0
pdays         -1.0
previous        0.0
dtype: float64
```

```
In [9]: bank.mean()
```

```
Out[9]: age            40.936210
balance        1362.272058
day            15.806419
duration       258.163080
campaign        2.763841
pdays         40.197828
previous        0.580323
dtype: float64
```

```
In [10]: bank.min()
```

```
Out[10]: age            18
job            admin.
marital        divorced
education      primary
default        no
balance        -8019
housing        no
loan           no
contact        cellular
day            1
month          apr
```

```

duration      0
campaign      1
pdays       -1
previous      0
poutcome     failure
y             no
dtype: object

```

```
In [11]: bank.max()
```

```


Out[11]: age      95
job      unknown
marital   single
education unknown
default   yes
balance  102127
housing   yes
loan      yes
contact   unknown
day       31
month     sep
duration  4918
campaign   63
pdays    871
previous  275
poutcome  unknown
y         yes
dtype: object

```


What simple models have you tried

```
In [12]: df = bank.copy()
```

```
In [13]: bank['job'].value_counts().plot(kind='bar');
```



```
In [14]: ax = sns.countplot(x = df["y"]) #Imbalanced dataset
plt.show()
```



categorical variables = ["job", "marital", "education", "default", "housing", "loan", "contact", "month", "day_of_week", "poutcome"]

numerical variables = ["age", "duration", "campaign", "pdays", "previous", "emp.var.rate", "cons.price.idx", "cons.conf.idx", "euribor3m", "nr.employed"]

```
In [15]: from sklearn.preprocessing import LabelEncoder, OrdinalEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
```

```
In [16]: objfeatures = df.select_dtypes(include="object").columns
le = LabelEncoder()

for feat in objfeatures:
    df[feat] = le.fit_transform(df[feat].astype(str))
```

```
In [17]: X = df.drop('y', 1)
y = df['y']

X = StandardScaler().fit_transform(X.astype(int))
```

```
In [18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4)
```

```
In [19]: from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree
from sklearn import tree
```

```
In [20]: short_tree = tree.DecisionTreeClassifier(max_depth = 3)
short_tree = short_tree.fit(X_train, y_train)
y_pred = short_tree.predict(X_test)
print('accuracy %2.2f ' % accuracy_score(y_test,y_pred))

cm = confusion_matrix(le.inverse_transform(y_test), le.inverse_transform(y_pred))
```

```
#test_results1 = pd.DataFrame(cm,index=Labels,columns=Labels)
#display(test_results1)#accuracy 0.66
```

accuracy 0.89

```
In [30]: y.value_counts(normalize=True)
```

```
Out[30]: 0    0.883015
         1    0.116985
         Name: y, dtype: float64
```

Since the data is imbalanced f1 score or auc is a better metric to compare models

```
In [34]: from sklearn.metrics import accuracy_score, f1_score, roc_auc_score
         depths = []
         accs = []
         trainaccuracy = []
         f1_train = []
         f1_test = []
         auc_train = []
         auc_test = []
         for i in range(3, 11):
             short_tree = tree.DecisionTreeClassifier(max_depth = i)
             short_tree = short_tree.fit(X_train, y_train)
             y_pred = short_tree.predict(X_test)
             y_prob = short_tree.predict_proba(X_test)[:,:1]
             acc = accuracy_score(y_test, y_pred)
             train_pred = short_tree.predict(X_train)
             train_prob = short_tree.predict_proba(X_train)[:,:1]
             auc_train.append(roc_auc_score(y_train, train_prob))
             auc_test.append(roc_auc_score(y_test, y_prob))
             f1_train.append(f1_score(y_train, train_pred))
             f1_test.append(f1_score(y_test, y_pred))
             trainacc = accuracy_score(y_train, train_pred)
             depths.append(i)
             accs.append(1-acc)
             trainaccuracy.append(1-trainacc)
```


```
In [37]: plt.figure(figsize=(16,5))
         plt.subplot(1,3,1)
         plt.plot(depths,trainaccuracy, label = "train")
         plt.plot(depths, accs, label = "test")
         plt.xlabel('Max Depth')
         plt.ylabel('Error')
         plt.legend()

         plt.subplot(1,3,2)
         plt.plot(depths,f1_train, label = "train")
         plt.plot(depths, f1_test, label = "test")
         plt.xlabel('Max Depth')
         plt.ylabel('F1-Score')
         plt.legend()

         plt.subplot(1,3,3)
```

```
plt.plot(depths,auc_train, label = "train")
plt.plot(depths, auc_test, label = "test")
plt.xlabel('Max Depth')
plt.ylabel('AUC Score')
plt.legend()

plt.show()
```

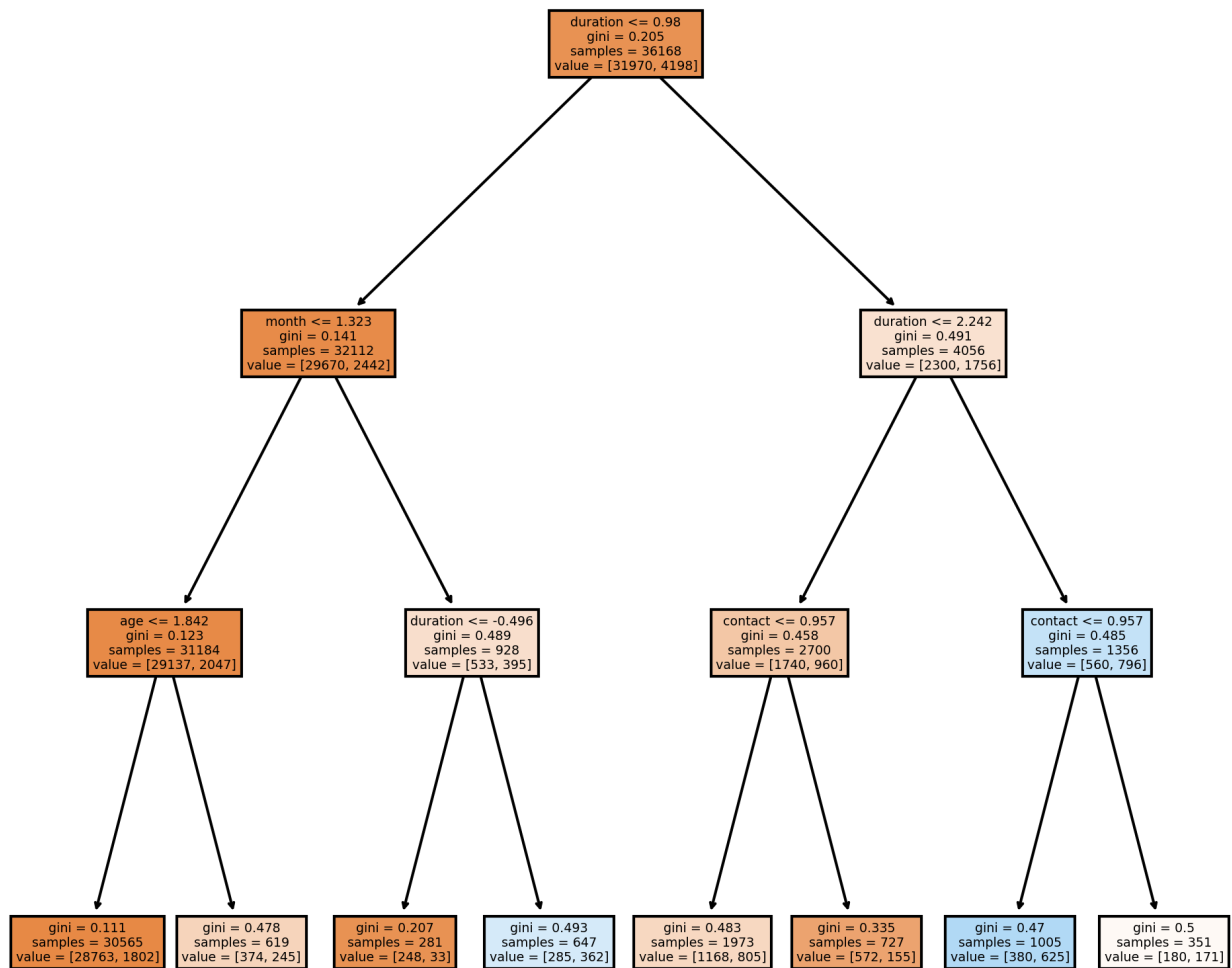


```
In [38]: pd.DataFrame({'Depth': depths, 'Train_Error': trainaccuracy, 'Test_Error': accs,
                      'Train_F1': f1_train, 'Test_F1': f1_test, 'Train_AUC': auc_train, 'Test_AU
                      #F1 score
```

```
Out[38]:
```

	Depth	Train_Error	Test_Error	Train_F1	Test_F1	Train_AUC	Test_AUC
0	3	0.107167	0.114121	0.337436	0.315650	0.744156	0.740511
1	4	0.104706	0.112352	0.360196	0.336815	0.805199	0.796087
2	5	0.099314	0.106491	0.496354	0.475204	0.845880	0.836646
3	6	0.093950	0.105385	0.514433	0.467895	0.869016	0.848791
4	7	0.091130	0.107044	0.528604	0.460424	0.892121	0.859422
5	8	0.084826	0.105607	0.570549	0.479564	0.913388	0.860801
6	9	0.079075	0.106823	0.591778	0.470975	0.923983	0.858337
7	10	0.072053	0.107597	0.639657	0.477163	0.936236	0.834246

```
In [45]: fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (8,8), dpi=300)
short_tree = tree.DecisionTreeClassifier(max_depth = 3)
short_tree = short_tree.fit(X_train, y_train)
tree.plot_tree(short_tree, filled=True, feature_names=df.drop(columns='y').columns);
```



Perceptron

```
In [46]: from sklearn.linear_model import Perceptron
clf = Perceptron(tol=1e-3, random_state=0)
clf.fit(X_train, y_train)
print("Acc:", clf.score(X_test, y_test))
```

Acc: 0.8458476169412805

```
In [61]: from sklearn.linear_model import LogisticRegression
```

```
In [72]: alphas = []
accs = []
trainaccuracy = []
f1_train = []
f1_test = []
auc_train = []
```

```

auc_test = []
for i in 10.0**np.arange(-7,2):
    clf = LogisticRegression(random_state=0,C=i,penalty='l2', )
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    y_prob = clf.predict_proba(X_test)[:,-1]
    acc = accuracy_score(y_test, y_pred)
    train_pred = clf.predict(X_train)
    train_prob = clf.predict_proba(X_train)[:,-1]
    auc_train.append(roc_auc_score(y_train, train_prob))
    auc_test.append(roc_auc_score(y_test, y_prob))
    f1_train.append(f1_score(y_train, train_pred))
    f1_test.append(f1_score(y_test, y_pred))
    trainacc = accuracy_score(y_train, train_pred)
    alphas.append(i)
    accs.append(1-acc)
    trainaccuracy.append(1-trainacc)

```

```

In [73]: pd.DataFrame({'alpha': alphas, 'Train_Error': trainaccuracy, 'Test_Error': accs,
                    'Train_F1': f1_train, 'Test_F1': f1_test, 'Train_AUC': auc_train, 'Test_AU

#F1 score

```

```

Out[73]:

```

	alpha	Train_Error	Test_Error	Train_F1	Test_F1	Train_AUC	Test_AUC
0	1.000000e-07	0.116069	0.120646	0.000000	0.000000	0.858413	0.854704
1	1.000000e-06	0.116069	0.120646	0.000000	0.000000	0.858614	0.854922
2	1.000000e-05	0.116069	0.120646	0.000000	0.000000	0.860421	0.856836
3	1.000000e-04	0.116152	0.120756	0.013155	0.003650	0.868168	0.865073
4	1.000000e-03	0.109959	0.114343	0.224303	0.229508	0.873432	0.871160
5	1.000000e-02	0.108079	0.112794	0.304077	0.301370	0.872555	0.870526
6	1.000000e-01	0.108300	0.111910	0.313409	0.318059	0.872050	0.870057
7	1.000000e+00	0.108300	0.112020	0.315090	0.318763	0.871985	0.869991
8	1.000000e+01	0.108245	0.112020	0.315679	0.318763	0.871979	0.869986

```

In [74]: plt.figure(figsize=(16,5))
plt.subplot(1,3,1)
plt.plot(alphas,trainaccuracy, label = "train")
plt.plot(alphas, accs, label = "test")
plt.xlabel('Max Depth')
plt.ylabel('Error')
plt.legend()
plt.xscale('log')


plt.subplot(1,3,2)
plt.plot(alphas,f1_train, label = "train")
plt.plot(alphas, f1_test, label = "test")
plt.xlabel('Max Depth')
plt.ylabel('F1-Score')
plt.legend()
plt.xscale('log')

```



```
plt.subplot(1,3,3)
plt.plot(alphas,auc_train, label = "train")
plt.plot(alphas, auc_test, label = "test")
plt.xlabel('Max Depth')
plt.ylabel('AUC Score')
plt.legend()
plt.xscale('log')

plt.show()
```



In []: